Assessment: Combining Tables
Question 1

1/1 point (graded)
You have created data frames `tab1` and `tab2` of state population and election data, similar to our module videos:

```
> tab1
state                  population
Alabama               4779736
Alaska                  710231
Arizona                 6392017
Delaware                897934
District of Columbia 601723

> tab2
state   electoral_votes
Alabama       9
Alaska        3
Arizona       11
California   55
Colorado      9
Connecticut   7

> dim(tab1)
[1] 5 2

> dim(tab2)
[1] 6 2
```

What are the dimensions of the table `dat`, created by the following command?

```
dat <- left_join(tab1, tab2, by = "state")
```

○ 3 rows by 3 columns

○ 5 rows by 2 columns

- ● 5 rows by 3 columns ✔

- ○ 6 rows by 3 columns

Answer

Correct:

When we use a left_join command, all rows in the left-hand table (in this case, tab1) are retained in the final table, so we expect to have five rows. In addition, columns from both tables will be included in the final "dat" table so we expect to have three columns.

| Submit | You have used 1 of 2 attempts |

ⓘ Answers are displayed within the problem

Question 2

1/1 point (graded)

We are still using the `tab1` and `tab2` tables shown in question 1. What join command would create a new table "dat" with three rows and two columns?

- ○ `dat <- right_join(tab1, tab2, by = "state")`

- ○ `dat <- full_join(tab1, tab2, by = "state")`

- ○ `dat <- inner_join(tab1, tab2, by = "state")`

- ● `dat <- semi_join(tab1, tab2, by = "state")` ✔

Answer

Correct:

The semi_join command takes tab1 and limits it to states that are also in tab2, without adding the additional columns in tab2. This gives us three rows (states in both tables) and two columns (state and population, the two columns in tab1).

| Submit | You have used 1 of 2 attempts |

ⓘ Answers are displayed within the problem

Question 3

1/1 point (graded)

Which of the following are real differences between the join and bind functions?

Please select all correct answers.

- ☑ Binding functions combine by position, while join functions match by variables. ✔

- ☑ Joining functions can join datasets of different dimensions, but the bind functions must match on the appropriate dimension (either same row or column numbers). ✔

- ☑ Bind functions can combine both vectors and dataframes, while join functions work for only for dataframes. ✔

- ☐ The join functions are a part of the dplyr package and have been optimized for speed, while the bind functions are inefficient base functions.

✔

| Submit | You have used 1 of 2 attempts |

ⓘ Answers are displayed within the problem

Question 4

1/1 point (graded)

We have two simple tables, shown below:

```
> df1
  x    y
  a    a
  b    a

> df2
  x    y
  a    a
  a    b
```

Which command would result in the following table?

```
> final
  x    y
  b    a
```

○ `final <- union(df1, df2)`

● ☑ `final <- setdiff(df1, df2)` ✔

○ `final <- setdiff(df2, df1)`

○ `final <- intersect(df1, df2)`

Answer
Correct: The `setdiff` command returns rows in df1 but not df2, which matches our table `final` .

Submit	You have used 1 of 2 attempts

ⓘ   Answers are displayed within the problem

Introduction to Questions 5-7

Install and load the `Lahman` library. This library contains a variety of datasets related to US professional baseball. We will use this library for the next few questions and will discuss it more extensively in the Regression course. For now, focus on wrangling the data rather than understanding the statistics.

The `Batting` data frame contains the offensive statistics for all baseball players over several seasons. Filter this data frame to define `top` as the top 10 home run (`HR`) hitters in 2016:

```
library(Lahman)
top <- Batting %>%
    filter(yearID == 2016) %>%
    arrange(desc(HR)) %>%      # arrange by descending HR count
    slice(1:10)      # take entries 1-10
top %>% as_tibble()
```

Also Inspect the `Master` data frame, which has demographic information for all players:

```
Master %>% as_tibble()
```

Question 5

1/1 point (graded)
Use the correct `join` or `bind` function to create a combined table of the names and statistics of the top 10 home run (HR) hitters for 2016. This table should have the player ID, first name, last name, and number of HR for the top 10 players. Name this data frame `top_names` .

Identify the `join` or `bind` that fills the blank in this code to create the correct table:

```
top_names <- top %>% _____ %>%
    select(playerID, nameFirst, nameLast, HR)
```

Which bind or join function fills the blank to generate the correct table?

- ○ `rbind(Master)`

- ○ `cbind(Master)`

- ● `left_join(Master)` ✔

- ○ `right_join(Master)`

- ○ `full_join(Master)`

- ○ `anti_join(Master)`

Answer code

```
top_names <- top %>% left_join(Master) %>%
    select(playerID, nameFirst, nameLast, HR)
top_names
```

[Submit]   You have used 1 of 2 attempts

ⓘ   Answers are displayed within the problem

Question 6

1/1 point (graded)

Inspect the `Salaries` data frame. Filter this data frame to the 2016 salaries, then use the correct bind join function to add a `salary` column to the `top_names` data frame from the previous question. Name the new data frame `top_salary`. Use this code framework:

```
top_salary <- Salaries %>% filter(yearID == 2016) %>%
    _____ %>%
  select(nameFirst, nameLast, teamID, HR, salary)
```

Which bind or join function fills the blank to generate the correct table?

○ `rbind(top_names)`

○ `cbind(top_names)`

○ `left_join(top_names)`

● `right_join(top_names)` ✔

○ `full_join(top_names)`

○ `anti_join(top_names)`

Answer code

```
top_salary <- Salaries %>% filter(yearID==2016) %>%
  right_join(top_names) %>%
  select(nameFirst, nameLast, teamID, HR, salary)
top_salary
```

| Submit | You have used 1 of 2 attempts |

ⓘ Answers are displayed within the problem

Question 7

2/2 points (graded)

Inspect the `AwardsPlayers` table. Filter awards to include only the year 2016.

How many players from the top 10 home run hitters won at least one award in 2016?
Use a set operator.

| 3 | ✔ **Answer:** 3 |

3

Answer code

```
Awards_2016 <- AwardsPlayers %>% filter(yearID == 2016)
length(intersect(Awards_2016$playerID, top_names$playerID))
```

How many players won an award in 2016 but were not one of the top 10 home run hitters in 2016?
Use a set operator.

44

✔ **Answer:** 44

44

Answer code

```
length(setdiff(Awards_2016$playerID, top_names$playerID))
```

Submit   You have used 1 of 10 attempts

ⓘ   Answers are displayed within the problem